Extending the stream Framework

Michael Hahsler
Southern Methodist University

Matthew Bolaños
Microsoft Corporation

John Forrest
Microsoft Corporation

Abstract

This document describes how to add new data stream sources DSD and data stream tasks DST to the **stream** framework.

Keywords: data streams, data mining, clustering.

1. Extending the stream framework

Since stream mining is a relatively young field and many advances are expected in the near future, the object oriented framework in **stream** is developed with easy extensibility in mind. Implementations for data streams (DSD) and data stream mining tasks (DST) can be easily added by implementing a small number of core functions. The actual implementation can be written in either R, Java, C/C++ or any other programming language which can be interfaced by R. In the following we discuss how to extend **stream** with new DSD and DST implementations.

1.1. Adding a new data stream source (DSD)

DSD objects can be a management layer on top of a real data stream, a wrapper for data stored in memory or on disk, or a generator which simulates a data stream with know properties for controlled experiments. Figure 1 shows the relationship (inheritance hierarchy) of the DSD classes as a UML class diagram (Fowler 2003). All DSD classes extend the abstract base class DSD. There are currently two types of DSD implementations, classes which implement R-based data streams (DSD_R) and MOA-based stream generators (DSD_MOA) provided in **streamMOA**. Note that abstract classes define interfaces and only implement common functionality. Only implementation classes can be used to create objects (instances). This mechanism is not enforced by S3, but is implemented in **stream** by providing for all abstract classes constructor functions which create an error.

The class hierarchy in Figure 1 is implemented using the S3 class system (Chambers and Hastie 1992). Class membership and the inheritance hierarchy is represented by a vector of class names stored as the object's class attribute. For example, an object of class DSD_Gaussians will have the class attribute vector c("DSD_Gaussians", "DSD_R", "DSD") indicating that the object is an R implementation of DSD. This allows the framework to implement all
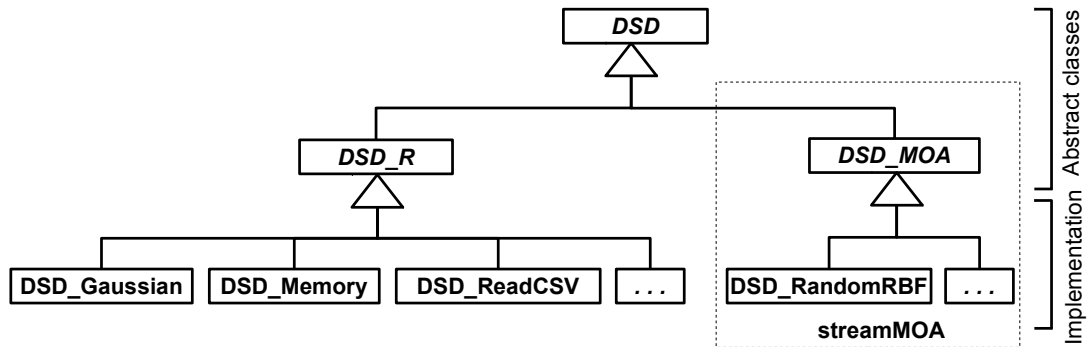
Figure 1: Overview of the data stream data (DSD) class structure.

common functionality as functions at the level of `DSD` and `DSD_R` and only a minimal set of functions is required to implement a new data stream source. Note that the class attribute has to contain a vector of all parent classes in the class diagram in bottom-up order.

For a new DSD implementation only the following two functions need to be implemented:

1. A creator function (with a name starting with the prefix `DSD_`) and

2. the `get_points()` method.

The creator function creates an object of the appropriate `DSD` subclass. Typically this S3 object contains a list of all parameters, an open R connection and/or an environment or a reference class for storing state information (e.g., the current position in the stream). Standard parameters are `d` and `k` for the number of dimensions of the created data and the true number of clusters, respectively. In addition an element called `"description"` should be provided. This element is used by `print()`.

The implemented `get_points()` needs to dispatch for the class and create as the output a data frame containing the new data points as rows. Also, if the ground truth (true cluster assignment as an integer vector; noise is represented by `NA`) is available, then this can be attached to the data frame as an attribute called `"cluster"`. If the new DSD implementation is capable of generating outliers, all outliers in the output data frame should be marked in a logical vector added as an attribute called `"outlier"`.

For a very simple example, we show here the implementation of `DSD_UniformNoise` available in the package's source code in file `DSD_UniformNoise.R`. This generator creates noise points uniformly distributed in a *d*-dimensional hypercube with a given range.

```
R> library("stream")

R> DSD_UniformNoise <- function(d = 2, range = NULL) {
+   if(is.null(range)) range <- matrix(c(0, 1), ncol = 2, nrow = d,
+     byrow = TRUE)
+   structure(list(description = "Uniform Noise Data Stream", d = d,
+     k = NA_integer_, range = range),
+         class = c("DSD_UniformNoise", "DSD_R", "DSD"))
```

Figure 2: Sample points from the newly implemented `DSD_UniformNoise` object.

```
+   }
R> get_points.DSD_UniformNoise <- function(x, n = 1,
+   assignment = FALSE, ...) {
+     data <- as.data.frame(t(replicate(n,
+       runif(x$d, min = x$range[ , 1], max = x$range[ , 2]))))
+     if(assignment) attr(data, "assignment") <- rep(NA_integer_, n)
+     data
+ }
```

The constructor only stores the description, the dimensionality and the range of the data. For this data generator k, the number of true clusters, is not applicable. Since all data is random, there is also no need to store a state. The `get_points()` implementation creates n random points and if assignments are needed attaches a vector with the appropriate number of `NA`s indicating that the data points are all noise.

Now the new stream type can already be used.

```
R> stream <- DSD_UniformNoise()
R> stream

Uniform Noise Data Stream
Class: DSD_UniformNoise, DSD_R, DSD
With NA clusters and NA outliers in 2 dimensions

R> plot(stream, main = description(stream))
```

The resulting plot is shown in Figure 2.

For the outlier data stream generator, we can take `DSD_Gaussians`. If we generate one cluster and one outlier for the horizon of 10 data points

```
R> stream <- DSD_Gaussians(k = 1, d = 2, outliers = 1, space_limit = c(0,0.5),
+                          outlier_options = list(outlier_horizon = 5))
```

we can obtain the first 10 data points, simultaneously looking for the cluster and outlier information. The data points obtained from the data stream are

```
R> points <- get_points(stream, n = 10, cluster = TRUE, outlier = TRUE)
R> points
```

```
       X1        X2
1   0.131   0.11701
2   0.335  -0.00470
3   0.211   0.20574
4   0.102   0.42236
5   0.225   0.16891
6   0.242   0.11392
7   0.282   0.09495
8   0.243   0.09585
9   0.145   0.00989
10  0.321   0.05669
```

we can extract cluster information as

```
R> attr(points, "cluster")
```

```
 [1] 1 1 1 2 1 1 1 1 1 1
```

and outlier marks as

```
R> attr(points, "outlier")
```

```
 [1] FALSE FALSE FALSE  TRUE FALSE FALSE FALSE FALSE FALSE FALSE
```

Several more complicated examples are available in the package's source code directory in files starting with DSD_.

1.2. Adding a new data stream tasks (DST)

DST refers to any data mining task that can be applied to data streams. The design is flexible enough for future extensions including even currently unknown tasks. Figure 3 shows the class hierarchy for DST. It is important to note that the DST base class is shown merely for conceptual purpose and is not directly visible in the code. The reason is that the actual implementations of data stream operators (DSO), clustering (DSC), classification (DSClass) or frequent pattern mining (DSFPM) are typically quite different and the benefit of sharing methods would be minimal.

DST classes implement mutable objects which can be changed without creating a copy. This is more efficient, since otherwise a new copy of all data structures used by the algorithm
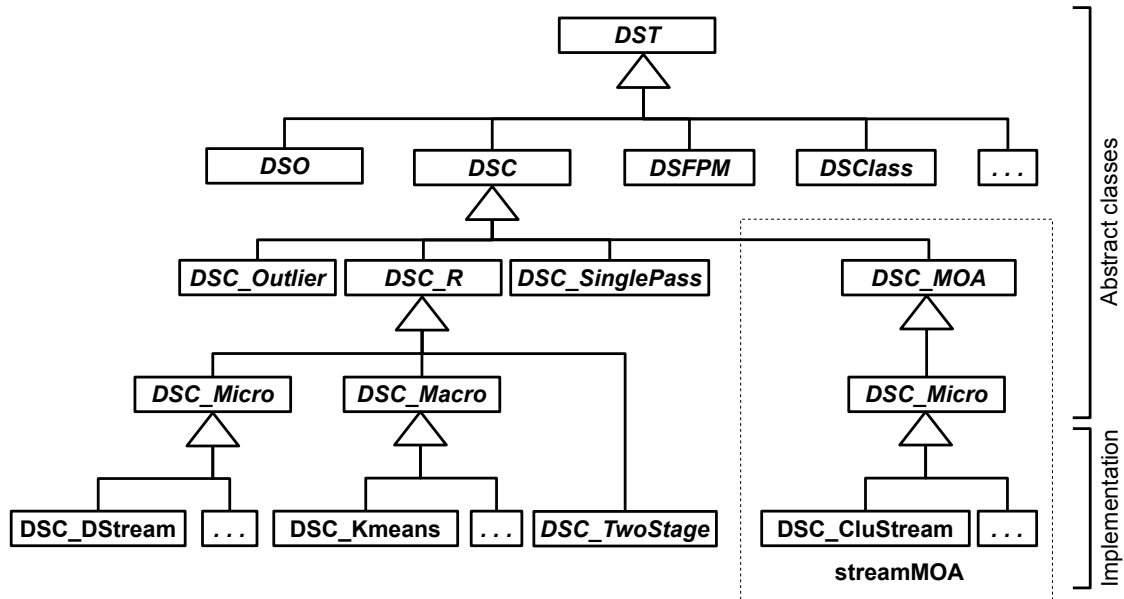
Figure 3: Overview of the data stream task (DST) class structure with subclasses for data stream operators (DSO), clustering (DSC), classification (DSClass) and frequent pattern mining (DSFPM).

would be created for processing each data point. Mutable objects can be implemented in R using environments or the recently introduced reference class construct (see package **methods** by the R Core Team (2014)). Alternatively, pointers to external data structures in Java or C/C++ can be used to create mutable objects.

To add a new data stream mining tasks (e.g., frequent pattern mining), a new package with a subclass hierarchy similar to the hierarchy in Figure 3 for data stream clustering (DSC) can be easily added. This new package can take full advantage of the already existing infrastructure in **stream**. An example is the package **streamMOA** Hahsler and Bolanos (2015), which can be used as a model to develop a new package. We plan to provide more add-on packages to **stream** for frequent pattern mining and data stream classification in the near future.

In the following we discuss how to interface an existing algorithm with **stream**. We concentrate again on clustering, but interfacing algorithms for other types of tasks is similar. To interface an existing clustering algorithm with **stream**,

1. a creator function (typically named after the algorithm and starting with DSC_) which created the clustering object,

2. an implementation of the actual cluster algorithm, and

3. accessors for the clustering

are needed. The implementation depends on the interface that is used. Currently an R interface is available as DSC_R and a MOA interface is implemented in DSC_MOA (in **streamMOA**).

The implementation for `DSC_MOA` takes care of all MOA-based clustering algorithms and we will concentrate here on the R interface.

For the R interface, the clustering class needs to contain the elements `"description"` and `"RObj"`. The description needs to contain a character string describing the algorithm. RObj is expected to be a reference class object and contain the following methods:

1. `cluster(newdata, ...)`, where `newdata` is a data frame with new data points.

2. `get_assignment(dsc, points, ...)`, where the clusterer `dsc` returns cluster assignments, outlier marks, and outlier identifiers for the input `points` data frame.

3. For micro-clusters: `get_microclusters(...)` and `get_microweights(...)`

4. For macro-clusters: `get_macroclusters(...)`, `get_macroweights` and `microToMacro(micro, ...)` which does micro- to macro-cluster matching.

5. For outlier detectors:

   - `clean_outliers(dsc, ...)` instructing the outlier detector to clean up the list of outliers

   - `get_outlier_positions(dsc, ...)` retrieving spatial positions of all current outliers

   - `recheck_outlier(dsc, outlier_id, ...)` re-checking the validity of the outlier by using its identifier, i.e., whether the outlier became an inlier in the meantime. This function must return `TRUE` if the outlier is still valid, and `FALSE` if the outlier has become an inlier in the meantime. Some outlier detectors allow outliers to decay (or fade), which rises an open question about whether a decayed outlier remains an outlier.

   - `noutliers(dsc, ...)` returns the number of current outliers.

Note that these are methods for reference classes and do not contain the called object in the parameter list. Neither of these methods are called directly by the user. Figure 4 shows that the function `update()` is used to cluster data points, and `get_centers()` and `get_weights()` are used to obtain the clustering. These user facing functions call internally the methods in RObj via the R interface in class `DSC_R`.

Single-pass clusterers and outlier detectors

Single-pass clusterers are processing input each data point separately. Processing is done in two steps. In the first step clusterer makes the classification and assessment. This classification is taken as the output result. In the second step, the clusterer makes necessary model updates using the input data point. Single-pass clusterers need to use the abstract class `DSC_SinglePass` anywhere between the abstract class `DSC` and the final clusterer class. For example:

```
R> DSC_MyClusterer <- function(x) {
+    structure(
+      list(
```

Figure 4: Interaction between the DSD and DSC classes.

```
+        description = "My new clusterer",
+        RObj = x
+    ), class = c("DSC_MyClusterer", "DSC_SinglePass", "DSC_Outlier",
+                 "DSC_Micro", "DSC_R", "DSC")
+  )
+ }
```

Figure 5 shows the interaction in case of single-pass clusterers. Obviously, since the model update is done at the end of processing for each data point, there is no need to perform `update()` before `get_assignment()`.

Outlier detectors are the clusterers that inherit the abstract class `DSC_Outlier`, placed anywhere between the abstract class `DSC` and the concrete final class, as seen in the previous code example. Besides all the method enumerated previously, outlier detectors must return additional structures from their `get_assignment()` method. For example:

```
R> stream <- DSD_Gaussians(k = 1, d = 2, outliers = 1,
+                          space_limit = c(0, 1), variance_limit = .01,
+                          outlier_options = list(outlier_horizon = 20))
R> points <- get_points(stream, n=20, cluster = TRUE, outlier = TRUE)
R> dsc <- DSC_MyClusterer()
R> assigns <- get_assignment(dsc, points, type="micro")
```

All outlier must have present their identifiers in the `attr(assigns, "outlier_corrid")`. Using these identifiers, calling the method `recheck_outlier(dsc, outlier_id, ...)` we can re-check the outlier validity.

For a comprehensive example of a clustering algorithm implemented in R, we refer the reader to `DSC_DStream` (in file `DSC_DStream.R`) in the package's R directory.
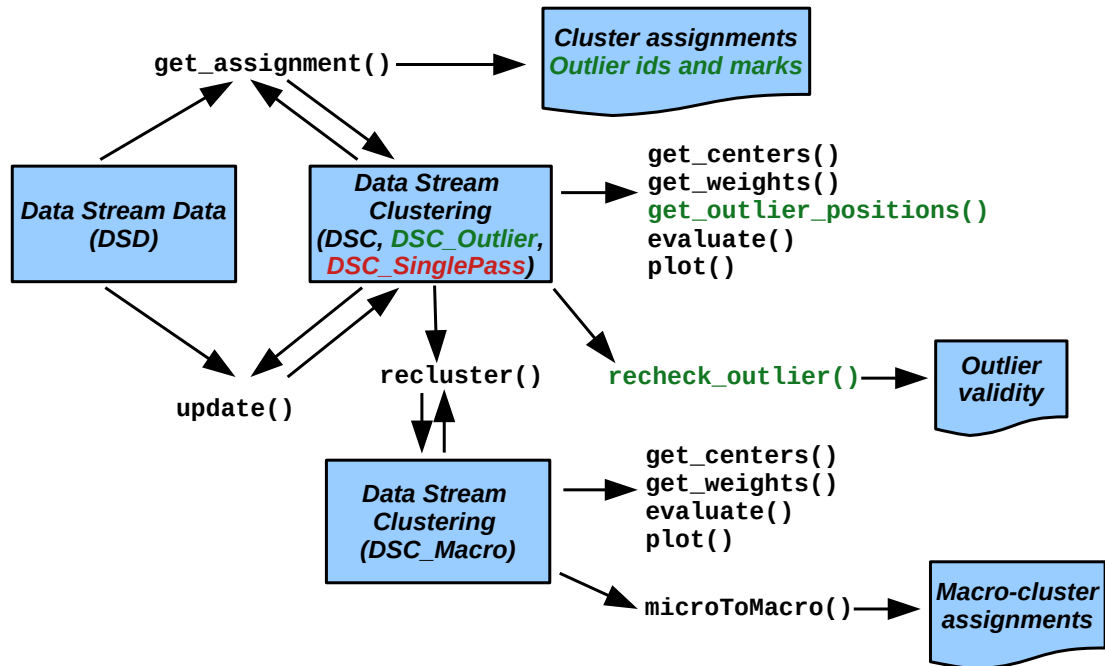
Figure 5: Interaction between the DSD and DSC classes for single-pass clusterers.

References

Chambers JM, Hastie TJ (1992). Statistical Models in S. Chapman & Hall. ISBN 9780412830402.

Fowler M (2003). UML Distilled: A Brief Guide to the Standard Object Modeling Language. 3 edition. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA. ISBN 0321193687.

Hahsler M, Bolanos M (2015). streamMOA: Interface for MOA Stream Clustering Algorithms. R package version 1.1-2, URL http://CRAN.R-project.org/package=streamMOA.

R Core Team (2014). R: A Language and Environment for Statistical Computing. R Foundation for Statistical Computing, Vienna, Austria. URL http://www.R-project.org/.

Affiliation:

Michael Hahsler
Engineering Management, Information, and Systems
Lyle School of Engineering
Southern Methodist University
P.O. Box 750122
Dallas, TX 75275-0122
E-mail: mhahsler@lyle.smu.edu
URL: http://lyle.smu.edu/~mhahsler

Matthew Bolaños
Research Now
5800 Tennyson Pkwy # 600
Plano, TX 75024 E-mail: mbolanos@curiouscrane.com

John Forrest
Microsoft Corporation
One Microsoft Way
Redmond, WA 98052-7329
E-mail: jforrest@microsoft.com